

Logistic Regression and Classification

Physical Address for Prof. Ilias Tagkopoulos

Computer Science:

Office: 3063 Kemper Hall

Phone: (530) 752-4821

Fax: (530) 752-4767

Instructor: Ilias Tagkopoulos
iliast@ucdavis.edu

Genome and Biomedical Sciences Facility:

Office: 5313 GBSF

Phone: (530) 752-7707

Fax: (530) 754-9658

■ From last lecture: Remember the definitions

- You were given a dataset with **m samples** $D = \{(x^{(i)}, y^{(i)}); i = 1 \dots m\}$.
 - Note that the superscript $x^{(i)}$ is the index of the sample.
 - Assume that each sample has **n attributes** (features)

$$D = \left\{ \begin{bmatrix} 1 & x_1^{(1)} & \dots & x_n^{(1)} \\ 1 & x_1^{(2)} & \dots & x_n^{(2)} \\ \dots & \dots & \dots & \dots \\ 1 & x_1^{(m)} & \dots & x_n^{(m)} \end{bmatrix} \begin{bmatrix} y^{(1)} \\ y^{(2)} \\ \dots \\ y^{(m)} \end{bmatrix} \right\} \text{ or}$$

$$D = \left\{ \begin{bmatrix} (x^{(1)})^T \\ (x^{(2)})^T \\ \dots \\ (x^{(m)})^T \end{bmatrix} \begin{bmatrix} y^{(1)} \\ y^{(2)} \\ \dots \\ y^{(m)} \end{bmatrix} \right\} \text{ with } x^{(i)} = \begin{bmatrix} 1 \\ x_1^{(i)} \\ \dots \\ x_n^{(i)} \end{bmatrix} \text{ being the } i^{th} \text{ sample}$$

$$D = \{X|Y\}$$

■ Previous lecture: Formulating a Regression problem

- By using the dataset D , you can **build a function that relates the input (x) to the output (y)**.
- You can then use this function to **predict** what the output (y) will be, based on a new, possibly never seen before, input.

$$f(x; w): X \rightarrow Y$$

With **weight vector** $w = \begin{bmatrix} w_0 \\ \cdots \\ w_n \end{bmatrix}$ and **input vector** $x^{(i)} = \begin{bmatrix} 1 \\ x_1^{(i)} \\ \cdots \\ x_n^{(i)} \end{bmatrix}$

- The task is to **find the optimal parameter values** for this function, i.e. the w , so that the function $f(x; w)$ ***best describes the relationship between X and Y***

■ Solutions for Linear Regression

- **Method 1: Ordinary Least Squares**

$$\frac{\partial RSS}{\partial \mathbf{w}} = \mathbf{0} \Rightarrow$$
$$\mathbf{w} = (X^T X)^{-1} X^T Y$$

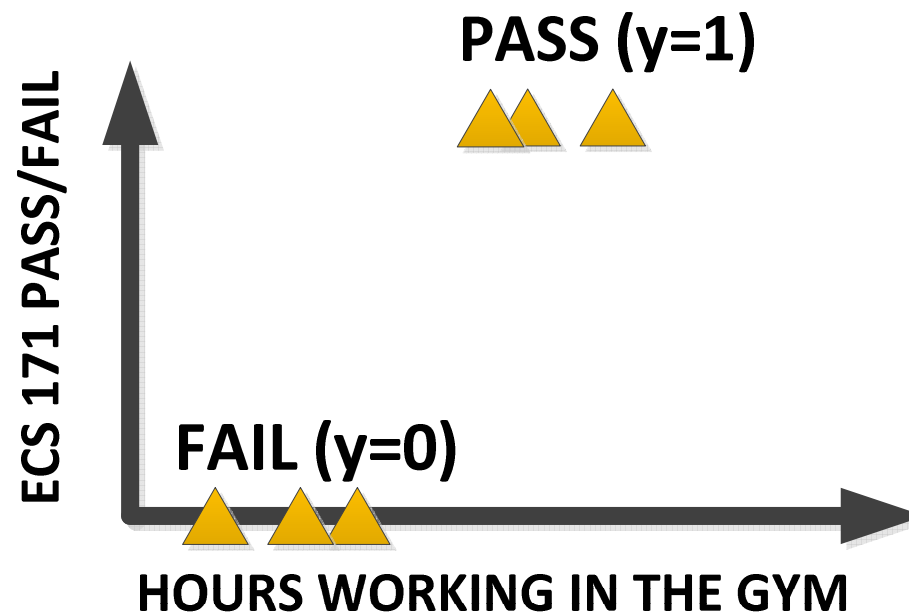
- **Method 2: Gradient Descent**

$$\boxed{w_j} := \boxed{w_j} + \boxed{a} \left(y^{(i)} - \sum_{k=0}^n w_k x_k^{(i)} \right) x_j^{(i)}$$

Next w_j Previous w_j constant Update proportional to error


■ Logistic Regression: a classification method

- Suppose that all your aim is not to find the value of a continuous variable *per se*, but to **categorize the samples into buckets or classes**.
- Sure, you can still perform linear regression and then threshold it, but the solution will be **sensitive to outliers** and to the **selected threshold**.




■ Logistic Regression: a classification method

- Suppose that all your aim is not to find the value of a continuous variable *per se*, but to **categorize the samples into buckets or classes**.
- Sure, you can still perform linear regression and then threshold it, but the solution will be **sensitive to outliers** and to the **selected threshold**.



■ Logistic Regression: a classification method


- Suppose that all your aim is not to find the value of a continuous variable *per se*, but to **categorize the samples into buckets or classes**.
- Sure, you can still perform linear regression and then threshold it, but the solution will be **sensitive to outliers** and to the **selected threshold**.



■ Logistic Regression: a classification method

- A better way to perform classification when within the regression framework is **logistic regression**
 - Same formulation as with linear regression, but instead of a polynomial, use the **sigmoid/logit/logistic function**:

$$\textit{sigm}(z) = \frac{1}{1 + e^{-z}}$$



■ Logistic Regression: a classification method

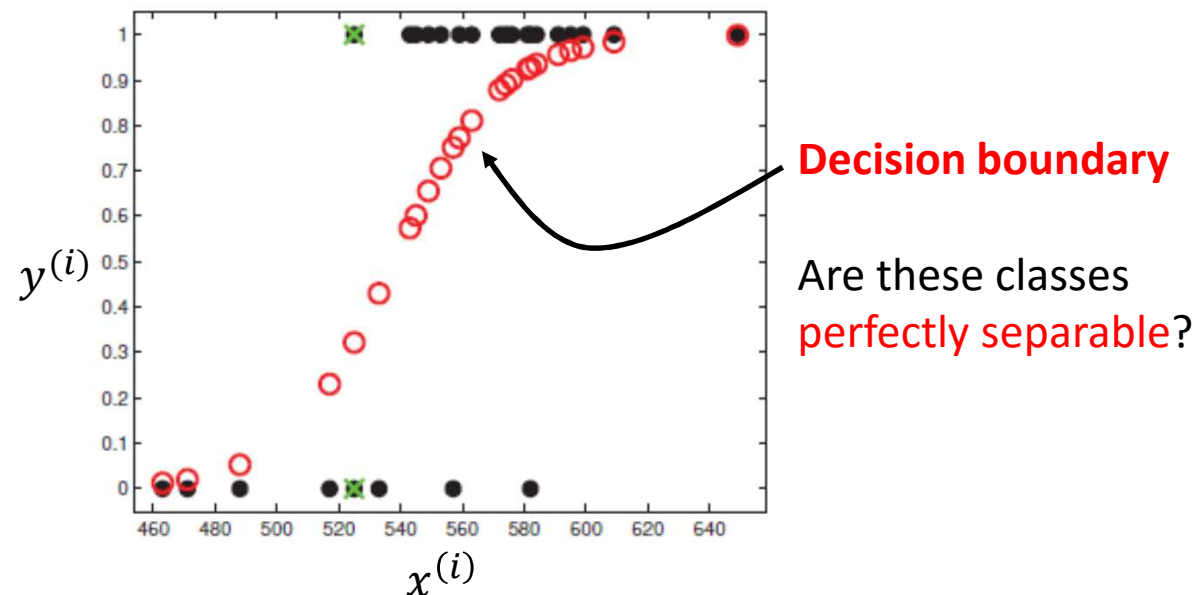
- As such, we can create the function

$$g(x; w) = \text{sigm}(w^T x) = \frac{1}{1 + e^{-w^T x}}$$

And categorize into **two classes (positive or negative)** for any given sample i , by defining the **label $y^{(i)}$** to be:

$$y^{(i)} = \begin{cases} 0 & \text{if } g(w^T x^{(i)}) < \text{threshold} \\ 1 & \text{if } g(w^T x^{(i)}) \geq \text{threshold} \end{cases}$$

With logistic regression, we can express it as **probabilities** too (instead of hard boundaries)



■ Logistic Regression: MLE

- Great, but how do we find the optimal parameter/weight \mathbf{w} set?
- There is no closed form solution (such as the OLS for linear regression) but we can **first formulate it as a MLE problem** and then **use gradient descent** to find the parameters.
- To do that, let's express the probability of each class:

$$\begin{aligned}P(y^{(i)} = 1 | x^{(i)}; \mathbf{w}) &= g(x^{(i)}; \mathbf{w}) \\P(y^{(i)} = 0 | x^{(i)}; \mathbf{w}) &= 1 - g(x^{(i)}; \mathbf{w})\end{aligned}$$

Which can also be written as:

$$p(y^{(i)} | x^{(i)}; \mathbf{w}) = g(x^{(i)}; \mathbf{w})^{y^{(i)}} (1 - g(x^{(i)}; \mathbf{w}))^{1-y^{(i)}}$$

■ Logistic Regression: MLE

- Assuming the samples are i.i.d. then we can write the log Likelihood as :

$$\begin{aligned} l(w) &\triangleq \log p(D|w) = \sum_{i=1}^M \log p(y^{(i)}|x^{(i)}; w) = \\ &= \sum_{i=1}^M \log \left(g(x^{(i)}; w)^{y^{(i)}} (1 - g(x^{(i)}; w))^{1-y^{(i)}} \right) = \\ &= \sum_{i=1}^M y^{(i)} \log g(x^{(i)}; w) + (1 - y^{(i)}) \log(1 - g(x^{(i)}; w)) \end{aligned}$$

- To maximize the log likelihood, we first find **the derivative of the log likelihood with respect to w:**

$$\frac{\partial l(w)}{\partial w}$$

Which yields

$$\frac{\partial l(w)}{\partial w_j} = \left(y^{(i)} - g(x^{(i)}; w) \right) x_j^{(i)}$$

■ Logistic Regression: Gradient Descent

- With that, we can now apply gradient descent (ascent), which as we know updates the w based on the following rule:

$$w_j := w_j + a \frac{\partial l(w)}{\partial w_j}$$

Or

$$w_j := w_j + a(y^{(i)} - g(x^{(i)}; w))x_j^{(i)}$$

Have you seen this before?

- Btw, gradient descent is not the only method to find the parameter w . E.g. Newton's method

$$w_j := w_j - \frac{\frac{\partial l(w)}{\partial w_j}}{\frac{\partial^2 l(w)}{(\partial w_j)^2}}$$


Logistic Regression: Newton-Raphson

Algorithm 8.1: Newton's method for minimizing a strictly convex function


```
1 Initialize  $\theta_0$ ;  
2 for  $k = 1, 2, \dots$  until convergence do  
3   Evaluate  $\mathbf{g}_k = \nabla f(\theta_k)$ ;  
4   Evaluate  $\mathbf{H}_k = \nabla^2 f(\theta_k)$ ;  
5   Solve  $\mathbf{H}_k \mathbf{d}_k = -\mathbf{g}_k$  for  $\mathbf{d}_k$ ;  
6   Use line search to find stepsize  $\eta_k$  along  $\mathbf{d}_k$ ;  
7    $\theta_{k+1} = \theta_k + \eta_k \mathbf{d}_k$ ;
```

$$w_j := w_j - \frac{\frac{\partial l(w)}{\partial w_j}}{\frac{\partial^2 l(w)}{(\partial w_j)^2}}$$

$$\theta_{k+1} = \theta_k - \eta_k \mathbf{H}_k^{-1} \mathbf{g}_k$$



(a)



(b)

■ Perceptron

- Actually if we use the same update rule but with hard boundaries, forcing the output to be $\{0,1\}$, we have the **perceptron learning algorithm**

$$w_j := w_j + a(y^{(i)} - g(x^{(i)}; w))x_j^{(i)}$$

with

$$g(z) = \begin{cases} 0 & \text{if } z < \text{threshold} \\ 1 & \text{if } z \geq \text{threshold} \end{cases}$$

Threshold can be any **scalar** (e.g. 0).

- Generally **Stochastic Gradient Descent** on **logistic regression** is faster and has a better performance.



End of Lecture 4