

Zhen Zhang

STA141 Assignment4 Code

```
load(url("http://eeyore.ucdavis.edu/stat141/Data/vehicles.rda"))
```

```
## Problem 1
```

```
# First I will introduce some functions I use to help me extract the pattern by gregexpr,  
# combine the regular expressions in multiple patterns, and give me the summary statistics  
# of the matched values.
```

```
extract_from_data <- function(pattern, original_string = vposts$body) {  
  # This function operates by first using gregexpr to extract the pattern I desire,  
  # and then calling extract_from_gregexpr to process the gregexpr result.
```

```
  grepexpr_index <- gregexpr(pattern, original_string, perl = T)  
  # make it invisible to the console  
  invisible(extract_from_gregexpr(grepexpr_index, original_string))  
}
```

```
extract_from_gregexpr <- function(grepexpr_index, original_string = vposts$body) {  
  # This function deals with the gregexpr result and extract the string implied by the  
  # starting index and the match.length attribute. I use the Map function to iterate  
  # the indices and the original string simultaneously, and at each row, extract the  
  # relevant information.
```

```
  Map(function(x, y) {  
    # how many matches in total  
    N <- length(x)  
    result <- sapply(1:N, function(i) {  
      # extract the string of each match  
      result_each <- substr(y, x[i], x[i] - 1 + attr(x, 'match.length')[i])  
      # if none is matched, set " to NA  
      ifelse(result_each == "", NA, result_each)  
    })  
  }, grepexpr_index, original_string)  
}
```

```
stat_for_total_num <- function(body) {  
  # This function calculates the statistics to remind me how many is matched.
```

```
# extract the Non-NA number and print it out
```

```

result <- sum(!is.na(body))
cat("\n")
cat(sprintf("The number of matching is %g", result))
cat("\n")
invisible(result)
}

```

```

paste_re <- function(x, y) {
  # This function combines different patterns, which is the main function renders to
  # Reduce to perform concatenation.

```

```

  paste(x, y, sep = "|")
}

```

```

# Before the formal regular expression matching, I would like to talk about how my matching
# process is organized. The first step is to combine different matching patterns by calling
# `paste_re` and `Reduce`. Then, I will use the `extract_from_data` function to extract
# pattern from string, usually `vposts$body` by default, but there are other cases. Since
# there may be multiple matching for each row, I will then select a single one by some
# algorithm. At last, I will show how many is matched using this pattern by
# `stat_for_total_num`

```

```

#### (a)

```

```

# Here I observe two patterns, one is for example, $6000 or 6000$, another is $6,000 or
# 6,000$. The first pattern is very easy, and to capture the second pattern, I use
# (?:pattern) to group digits and the corresponding ','.
price_pattern <- c("\\$[^0-9]{0,4}(?:[0-9]+[,])*[0-9]*", "(?:[0-9]+[,])*[0-9]*\\$[^0-9]{0,4}")

```

```

# extract relevant pattern strings

```

```

price_from_body <- extract_from_data(Reduce(paste_re, price_pattern))

```

```

# The first step is to remove $ symbol, and the second step is to remove ',' symbol if
# possible, and the last step is for multiple prices, return the largest one.

```

```

price_from_body_single <- sapply(price_from_body, function(x) {
  N <- length(x)
  price_total <- sapply(1:N, function(i) {
    # remove $
    price <- gsub("[^0-9]*([0-9,]+).*", "\\1", x[i])
    # remove ,
    price <- gsub("[,]", "", price)
    # sometimes the price and the year is concatenated together, so I remove them
    price <- ifelse(nchar(price) > 8, substr(price, 1, nchar(price) - 4), price)
  })
  # return the max price
  max(as.numeric(price_total))
}

```

```
})
```

```
# show how many is matched successfully.
```

```
stat_for_total_num(price_from_body_single)
```

```
# show how many is consistent with the price column
```

```
sum(price_from_body_single == vposts$price, na.rm = T)
```

```
# show some results I get
```

```
head(price_from_body_single, n = 10)
```

```
### (b)
```

```
# According to the VIN standard, VIN has 17 characters, where there is no latter I, O, Q.
```

```
# In addition, the last six letters must be digits, so my regular expression is written
```

```
# as follows:
```

```
VIN_pattern <- c('(![0-9]{17})[0-9A-HJ-NPR-Z]{13}[0-9]{4}')
```

```
# There is only one pattern, so I will not use Reduce function.
```

```
VIN_from_body <- extract_from_data(VIN_pattern)
```

```
# Here I will select the VIN if there are multiple results, since the VIN usually comes
```

```
# at first position of the body text.
```

```
VIN_from_body_single <- sapply(VIN_from_body, '[', 1)
```

```
# show how many is matched successfully.
```

```
stat_for_total_num(VIN_from_body_single)
```

```
# add to vposts
```

```
vposts$VIN <- VIN_from_body_single
```

```
# show some results I get
```

```
head(VIN_from_body_single, n = 10)
```

```
### (c)
```

```
# At first it may seem very easy to extract the pattern of price. But in the last line,
```

```
# there is '(916) 715-31 SEVEN ZERO'. It will not be extracted by the normal expression
```

```
# only focusing on the digits. So at first I write a function to transform natural
```

```
# language to digits.
```

```
# This is the vector that combines the natural language digits
```

```
nums <- c('zero', 'one', 'two', 'three', 'four', 'five', 'six', 'seven', 'eight', 'nine')
```

```
num_to_digit <- function(text) {
```

```
# This is the function that transform natural language to digits. Note here my pattern
```

```
# is paste0('[ -]?', (nums[i]), '[ -]?'), since when there is natural languages, people
```

*# tend to separate them by space or hyphens. For example, (916) 715-31 SEVEN ZERO. This
pattern will remove the unnecessary spaces or hyphens after transforming to digits.*

```
for (i in seq_along(nums)) {  
  text <- gsub(paste0('[ -]?', (nums[i]), '[ -]?'), i - 1, text, ignore.case = T)  
}
```

return value

```
text  
}
```

call the function above, to return the new texts for body particular for phones.

```
body_modified_for_phone <- sapply(vposts$body, num_to_digit)
```

Now the string '(916) 715-31 SEVEN ZERO' can be transformed to:

```
num_to_digit('(916) 715-31 SEVEN ZERO')
```

Ok, it is fine now.

The pattern for phone number

*# there are three parts for a phone number, each part inside cannot be separated, and
between each part, there can be at most one space [] or hyphen [-]. Also, the first
part can be surrounded by parenthesis ().*

```
phone_pattern <- c('\\(?:[0-9]{3}\\)?[. -]*[0-9]{3}[. -]*[0-9]{4}')
```

extrac the phone numbers from body_modified_for_phone

```
phone_from_body <- extract_from_data(phone_pattern, body_modified_for_phone)
```

get the first phone number if there is multiple

```
phone_from_body_single <- sapply(phone_from_body, '[', 1)
```

show how many is matched successfully

```
stat_for_total_num(phone_from_body_single)
```

add to vposts

```
vposts$phone <- phone_from_body_single
```

show some results I get

```
head(phone_from_body_single, n = 10)
```

(d)

*# The email address always have @ symbol and xxx.xxx.com/net/org pattern. The thing should
be kept in mind is the subdomain, so I use (?:) to group characters and dot, and let
them emerge together. Also, keep com|org|net from emerging in the subdomain, since
they should only come into being in top domain name.*

```
email_pattern <- c('[:alnum:]]+@(?:?!com|org|net)[:alnum:]]+\\.)+(?:com|org|net)')
```

```

email_from_body <- extract_from_data(Reduce(paste_re, email_pattern))

# Here I will select the email if there are multiple results, since most of them are
# identical.
email_from_body_single <- sapply(email_from_body, '[', 1)

# show how many is matched successfully
stat_for_total_num(email_from_body_single)

# add to vposts
vposts$email <- email_from_body_single

# show some results I get
email_from_body_single[which(!is.na(email_from_body_single))]

#### (e)

# year pattern, including 19xx, 200x, 201x, 9x, 0x
year_pattern <- c('19[5-9][0-9]', '20[0-1][0-9]', '[09][0-9]')

# extract year by year_pattern
year_from_description <- extract_from_data(Reduce(paste_re, year_pattern), vposts$description)

# If there are multiple results, select the first one. Also, if there is 9x, 0x pattern,
# convert them to 199x, 200x.
year_from_description_single <- sapply(year_from_description, function(x) {
  result <- x[1]
  if (nchar(result) == 2 && !is.na(result)) {
    x_first_digit <- substr(result, 1, 1)
    if (x_first_digit == '0') {
      result <- paste0('20', result)
    } else {
      result <- paste0('19', result)
    }
  }
  result
})

# show how many is matched successfully.
stat_for_total_num(year_from_description_single)

# how many is wrong
sum(vposts$year != year_from_description_single, na.rm = T)

# show some results I get

```

```
head(year_from_description_single, n = 10)
```

```
### (f)
```

```
# Here I extract model name from two string texts, one is vposts$header, another is  
# vposts$description. Most of the model names come at the second word after time. Sometimes  
# there is no such string in vposts$header, so I will try to search it in  
# vposts$description.
```

```
# extract the maker name
```

```
vposts_maker_names <- unique(vposts$maker)
```

```
vposts_maker_names <- vposts_maker_names[!is.na(vposts_maker_names)]
```

```
time_identifier_model <- function(x) {
```

```
# This is the function that implements the first algorithm, time identifier. Using the  
# splitted result, it will first match the time pattern, after that, get the model index,  
# then get the model
```

```
# match the time pattern
```

```
x_time_index <- which(sapply('[0-9]{4}', grepl, x))
```

```
# get the model index
```

```
x_model_index <- x_time_index + 2
```

```
# get the model
```

```
x_model <- x[x_model_index]
```

```
# set unmatched to NA
```

```
if (length(x_model) == 0) {
```

```
  x_model <- NA
```

```
}
```

```
# return the result
```

```
x_model
```

```
}
```

```
vposts_split <- function(string) {
```

```
# This function split the string by identifier not digits, alphabets or hyphens. Most  
# of the time, the split one will be space, and there are some situations that , will  
# be the separate tag. Here I preserve -, since there are many vehicles that has the  
# model name of '3-Series'
```

```
string_elements <- unlist(strsplit(string, "[^0-9a-zA-Z-]", perl = T))
```

```
string_elements <- string_elements[string_elements != ""]
```

```
}
```

```
model_extract <- function(string) {
```

```
# This function first does some preparation, namely the split job, then render it to
```

```
# time identifier and maker identifier. It connects different small jobs into a whole  
# process.
```

```
unlist(lapply(string, function(x) {  
  # split the string  
  x_elements <- vposts_split(x)  
  
  # the first algorithm, using the time identifier  
  x_model <- time_identifier_model(x_elements)[1]  
}))  
}
```

```
model_extract_total <- function() {  
  # This function calculates the model extracted from header and description, and then  
  # will use the result from header by default, and only use the result from description  
  # if the result from header is NA.
```

```
  model_from_header <- model_extract(vposts$header)  
  model_from_description <- model_extract(vposts$description)
```

```
  Map(function(x, y) {  
    ifelse(!is.na(x), x, y)  
  }, model_from_header, model_from_description)  
}
```

```
model_from_header_description <- unlist(model_extract_total())
```

```
# convert to lower case, being case insensitive  
model_from_header_description <- tolower(model_from_header_description)
```

```
# show some results I get  
head(model_from_header_description, n = 10)
```

```
# *Correct the typos in model name*
```

```
# The algorithm is a little complicated, so I will explain the basic idea. First,  
# I will create a dataframe of maker to model, and then classify the data based the  
# maker. After that, for each maker, I will split the model by its frequency. If the  
# frequency is higher than a threshold (which is a parameter, confidence_threshold),  
# then it will be normal names, otherwise, I think it is mistyped. The mistyped names  
# will then calculate the distance with the normal names, see if it is close enough to  
# anyone of the normal names. Note here that the inclusion of partial = T when calculating  
# adist will give us also the result of abbreviations. The distance threshold is 1, so if  
# it is small than 1, the mistyped names will be substitute by the normal name, and if  
# there is a tie, it will be substituted by the most frequent normal name. Also, if none of  
# the distance is smaller than 1, then I also think it is not mistyped. At last, for each
```

*# mistyped names, I will create a rule controlling how it should be modified, which is the
result of this function.*

```
x_lower_names_transform <- function(x_lower_names, x_higher_names, x_adist, x_table) {  
  # This function change the mistyped names to normal names, by the adist matrix.  
  # The distance threshold is 1, so the distance is small than 1, the mistyped names will  
  # be substitute by the normal name, and if there is a tie, it will be substituted by  
  # the most frequent normal name. Also, if none of the distance is smaller than 1,  
  # then I also think it is not mistyped.
```

```
sapply(1:length(x_lower_names), function(i) {  
  # initialized the result  
  result <- x_lower_names[i]  
  # calculate the nearest distance  
  name_adist <- x_adist[i, ]  
  name_adist_min <- min(name_adist)  
  # if the distance is one, set the name to it  
  if (name_adist_min == 1) {  
    name_adist_min_index <- which(name_adist == name_adist_min)  
    name_adist_min_index_names <- x_higher_names[name_adist_min_index]  
    if (length(name_adist_min_index_names) == 1) {  
      result <- name_adist_min_index_names  
    } else {  
      # in case of tie, select the one with the most frequency  
      compare_table <- x_table[name_adist_min_index_names]  
      result <- names(which.max(compare_table))  
    }  
  }  
  # return value  
  result  
})  
}
```

```
x_change_chain_rule <- function(x_lower_names, x_lower_names_transformed) {  
  # create the change rule, and only return the one that is really changed under  
  # my algorithm, others will leave them unchanged, meaning the rule is NULL.  
  # Here I find the Map function is most convenient.
```

```
  Map(function(x, y) {  
    if (x != y) {  
      y  
    }  
  }, x_lower_names, x_lower_names_transformed)  
}
```

```
modify_typo_model_maker <- function(x, confidence_threshold) {
```



```

# This function operates on the abstract level of each maker, calculateing each level's
# lowest names, for each maker, and also convert the corresponding level to the nearest
# model. It first does some preparation, and then call x_lower_names_transform function
# to return the modify model values, and at last, call the x_change_chain_rule function
# to create the change chain.

# Some maker only has one model, so at that case there is nothing to do
if (length(x) > 1) {

  # calculate lowest names, higher names, table for future calculation
  x_table <- table(x)
  x_table_names <- names(x_table)
  x_lower_names <- x_table_names[x_table <= confidence_threshold]
  x_higher_names <- x_table_names[x_table > confidence_threshold]
  # calculate the adist matrix
  # It should be noted here that, using the option partial = T will also caputre
  # abbreviation cases.
  x_adist <- adist(x_lower_names, x_higher_names, partial = T)

  # I must ensure there is both the mistyped models and normal models
  if (length(x_adist) != 0) {

    # modify the model name, if the adist value is 1. if tie, choose the one with the
    # most frequency. If no adist value is 1 related to that name, keep it unchanged
    x_lower_names_transformed <- x_lower_names_transform(x_lower_names,
x_higher_names,
x_adist, x_table)

    # create the change rule
    x_change_chain <- x_change_chain_rule(x_lower_names, x_lower_names_transformed)

    # return the change rule
    x_change_chain
  }
}

modify_typo_model <- function(model_from_header, confidence_threshold) {
  # This function is the main function to implement the correct typos. It first calculates
  # the dataframe and the splitted dataframe by maker, then supplies these results to
  # modify_typo_model_maker function. It connects different parts of the correcting job,
  # and return the value I really want.

  # preparation

  # combine the data into a dataframe

```

```

maker_model_data <- data.frame(maker = vposts$maker, model = model_from_header,
                               stringsAsFactors = F)
# split by the maker
maker_model_data_split <- split(maker_model_data$model, maker_model_data$maker)

# Now comes to the main part. It operates on the abstract level of each maker.
# It calculate each level's lowest names, for each maker, and also convert to the
# nearest model
lapply(maker_model_data_split, modify_typo_model_maker, confidence_threshold)
}

# call the modify function to modify the model name
model_modify_rule <- modify_typo_model(model_from_header_description, 3)

# show the rule, with maker toyota, here if the value under a model is null, there is no
# modification to it, if there is some value under a model, this is the one should be
# modified to.
unlist(model_modify_rule['toyota'])

model_final <- unlist(lapply(1:length(vposts$header), function(i) {
  # present the values of maker, model, and modified model
  maker <- vposts$maker[i]
  model <- model_from_header_description[i]
  modified_model <- model_modify_rule[[maker]][[model]]
  # if it is not null, then there should be a rule to change it
  if (!is.null(modified_model)) model <- modified_model
  # return result
  model
})))

# add the model to the vposts dataframe
vposts$model <- model_final

# show some results
head(model_final, n = 10)

## Question two

# First I sort the table, and see which model is the most:
invisible(sort(table(model_final)))

# For the top five models, I want to see which maker they belong to:
top_five_models_names <- names(table(model_final))[
  order(table(model_final), decreasing = T)[1:5]]
# check if the model belong to the same maker
sapply(top_five_models_names, function(x) {

```

```

x_index <- which(model_final == x)
unique(vposts$maker[x_index])
})

```

I will pick the top two combination honda-civic, toyota-camry.

select civic indices and camry indices:

```

select_indices <- function(model) {
  # this is the function that can return the indices related to this model
  name_index <- which(model_final == model)
}

```

```

get_model_data <- function(model, intereste_cols_flag = TRUE) {
  # this funciton first calls select_indices, then use the indices to select rows
  # from vposts

```

```

  indices <- select_indices(model)
  # get the dataframe for price, age, odometer and condition
  model_data <- vposts[indices, ]
  model_data$age <- 2015 - model_data$year
  if (intereste_cols_flag) {
    model_data <- model_data[c('price', 'odometer', 'condition', 'age')]
  } else {
    model_data <- model_data[c('price', 'odometer', 'condition', 'age', 'city')]
  }
}

```

get complete data without NA

```

model_data <- model_data[complete.cases(model_data),]

```

since there is a factor condition, if the frequency is so low, it will not perform

cross-validation. So I remove such levels, with frequency less than 5.

```

model_data$condition <- factor(model_data$condition)
model_data_cond_table <- table(model_data$condition)
model_data_cond_names_high <- names(model_data_cond_table)[model_data_cond_table > 5]
model_data <- subset(model_data, condition %in% model_data_cond_names_high)

```

return

```

model_data
}

```

My Approach

Before All, I will first draw a map that prints different values of price, odometer, age and condition.

The basic idea of my implementation in this question is, first I will choose which model

*# to select. Since there are multiple models, I will only try some of them, including
linear regression (lm), knn, regression tree (tree). The selection method is to use the
package of caret to do cross validation on each of them, and see which one has the
lowest RMSE. After we have already got one model in hand, we still need to select
the parameters. For example, knn has the k parameter indicating how many nearest points
should be used to predict the point we are interested. So I will again use cross
validation to select the parameter, and at last, get the final model.*

Note that the assumptions here:

*# lm: The error terms are normally distributed around 0 with the same standard deviation,
identically independent distributed (i.i.d.).*

*# knn and regression tree: They are non-parametric methods, so there is no assumption about
them.*

load caret package
library('caret')

Part I, Model of civic

get the data
civic_data_total <- get_model_data('civic', FALSE)
civic_data <- get_model_data('civic')

show the data
head(civic_data, n = 10)

Step 0, draw the boxplot of city vs price.

library(ggplot2)
ggplot(civic_data_total, aes(city, price)) +
 geom_boxplot() +
 ggtitle("price vs city plot for model civic")

*# The result indicates that, city has an effect on the price of vehicles. While vehicles
in lasvegas and sac have high price, the vehicles in boston, chicago, denver, nyc and
sfbay have low price.*

Step I, find which model is the best

*# Before I really do the regression, I will first set the cross validation parameter.
The parameters I choose is 5-fold cross validation, repeating 3 times.*
set.seed(1234)
fitControl <- trainControl(method = 'cv', number = 5, repeats = 3)

The lm approach

```
civic_fit_lm <- train(price ~ ., data = civic_data, method = 'lm',  
  trControl = fitControl)  
civic_fit_lm
```

It is easily seen that the lm fit result is RMSE: 2427.636

The knn approach

```
civic_fit_knn <- train(price ~ ., data = civic_data, method = 'knn',  
  trControl = fitControl, tuneGrid = expand.grid(k = 5:20))  
civic_fit_knn  
plot(civic_fit_knn, main = 'plot for knn under model of civic')
```

The best model for knn is k = 15, with RMSE: 3042.217

The regression tree (tree) approach

```
civic_fit_tree <- train(price ~ ., data = civic_data, method = 'rpart',  
  trControl = fitControl, tuneGrid = expand.grid(maxdepth = 1:8))  
civic_fit_tree  
plot(civic_fit_tree, main = 'plot for regression tree under model of civic')
```

The best model for regression tree is k = 4, with RMSE: 1931.201

Now we have the best model, which is regression tree, so I will dig deeply in regression tree

Step II, find the best regression tree model

```
library(tree)
```

regression tree details

```
civic_fit_tree <- tree(price ~ ., data = civic_data)  
civic_fit_tree
```

plot the result

```
plot(civic_fit_tree)  
text(civic_fit_tree, pretty = 0)
```

use cross validation to confirm the best tree depth

```
civic_fit_tree_cv <- cv.tree(civic_fit_tree)  
civic_fit_tree_cv
```

plot the result, the relationship between nodes and deviation

```
plot(civic_fit_tree_cv$size, civic_fit_tree_cv$dev, type="b", xlab = 'size',  
  ylab = 'deviation', main = 'deviation vs terminal node size plot')
```

It seems size 5 is the best option

```
civic_fit_tree_prune <- prune.tree(civic_fit_tree, best = 5)
```

```
civic_fit_tree_prune
# plot the best result
plot(civic_fit_tree_prune)
text(civic_fit_tree, pretty = 0)
```

Now let's do a prediction, the data I use is the first observation:

```
civic_data[1,]
# let's use tree model to predict
civic_prediction_one <- predict(civic_fit_tree_prune, civic_data[1,])
civic_prediction_one
# It is very close!
```

Step III, Analyze the result

Let's look back the regression tree here:

```
civic_fit_tree_prune
```

*# It is seen that, only age is valid in determining price, which means price has little
relationship with odometer and condition. It may seem unreasonable at first, but consider
this: A vehicle with higher age usually comes with high odometer and poor condition.*

*# In other words, age is related with odometer and condition. So it is not wired that
price only relates with age.*

Part II, Model of camry

get the data

```
camry_data_total <- get_model_data('camry', FALSE)
camry_data <- get_model_data('camry')
```

show the data

```
head(camry_data, n = 10)
```

Step 0, draw the boxplot of city vs price.

```
library(ggplot2)
ggplot(camry_data_total, aes(city, price)) +
  geom_boxplot() +
  ggtitle("price vs city plot for model camry")
```

*# The result indicates that, city has almost no effect on price. The mean is almost
identical, and the only difference relates to the standard deviation, namely the
fluctuation of the price. The variance is high in city of boston, lasvegas and sac.*

Step I, find which model is the best

```

# As in part I, I will first set the cross validation parameter.
# The parameters I choose is also 5-fold cross validation, repeating 3 times.
set.seed(1234)
fitControl <- trainControl(method = 'cv', number = 5, repeats = 3)

# The lm approach
camry_fit_lm <- train(price ~ ., data = camry_data, method = 'lm',
                     trControl = fitControl)
camry_fit_lm

# It is easily seen that the lm fit result is RMSE: 2597.181

# The knn approach
camry_fit_knn <- train(price ~ ., data = camry_data, method = 'knn',
                     trControl = fitControl, tuneGrid = expand.grid(k = 5:20))
camry_fit_knn
plot(camry_fit_knn, main = 'plot for knn under model of camry')

# The best model for knn is k = 10, with RMSE: 3227.618

# The regression tree (tree) approach

camry_fit_tree <- train(price ~ ., data = camry_data, method = 'rpart2',
                     trControl = fitControl, tuneGrid = expand.grid(maxdepth = 1:8))
camry_fit_tree
plot(camry_fit_tree, main = 'plot for regression tree under model of camry')

# The best model for regression tree is k = 5, with RMSE: 2215.835

# Although knn is better than tree model, I will still use tree model, since it is more
# reasonable than a knn model with high k and easier to interpret.

# Step II, find the best regression tree model

library(tree)

# regression tree details
camry_fit_tree <- tree(price ~ ., data = camry_data)
camry_fit_tree
# plot the result
plot(camry_fit_tree)
text(camry_fit_tree, pretty = 0)

# use cross validation to confirm the best tree depth
camry_fit_tree_cv <- cv.tree(camry_fit_tree)
camry_fit_tree_cv

```

```
# plot the result, the relationship between nodes and deviation
plot(camry_fit_tree_cv$size, camry_fit_tree_cv$dev, type="b", xlab = 'size',
     ylab = 'deviation', main = 'deviation vs terminal node size plot')
```

```
# It seems size 5 is the best option
```

```
camry_fit_tree_prune <- prune.tree(camry_fit_tree, best = 5)
```

```
camry_fit_tree_prune
```

```
# plot the best result
```

```
plot(camry_fit_tree_prune)
```

```
text(camry_fit_tree, pretty = 0)
```

```
# Now let's do a prediction, the data I use is the first observation:
```

```
camry_data[1,]
```

```
# let's use tree model to predict
```

```
camry_prediction_one <- predict(camry_fit_tree_prune, camry_data[1,])
```

```
camry_prediction_one
```

```
# It is very close, again!
```

```
# Step III, Analyze the result
```

```
# Let's look back the regression tree here:
```

```
camry_fit_tree_prune
```

```
# It is quite like the situation of civic, so I will copy my previous analysis here:
```

```
# It is seen that, only age is valid in determining price, which means price has little
```

```
# relationship with odometer and condition. It may seem unreasonable at first, but consider
```

```
# this: A vehicle with higher age usually comes with high odometer and poor condition.
```

```
# In other words, age is related with odometer and condition. So it is not wired that
```

```
# price only relates with age.
```

```
# Conclusion
```

```
# To conclude, I will use my regression model, since I think the result is good, and
```

```
# it is reasonable, and can give me some guide.
```